

Exercise 1 - Canny Edge Detector

The goal of this exercise is to familiarise yourself with edge detection. You will implement the *Canny edge detector*, one of the most common edge detectors. You can see an example of the output of the algorithm in Figure 1. Your results will look slightly different depending on the choice of parameters and implementation details.

A code construct and an example image is provided in a zip file which can be downloaded [here](#). Use these files as templates for your program and refer to the comments in the files for details on expected inputs and outputs.

To setup your computer for the exercises, we provide an explanation on how to setup the Conda environment [here](#). You can also find a video going through the setup process as well as a Python crash course video on TUWEL.



Figure 1: Resulting edge image (left), original image overlaid with detected edges (right).

Part 1: Implementation (12 Points)

You should write code so that it can be understood by someone else (the tutors, for example). Comment sections and lines which might not be easy to understand. If you've thought about a line or section for more than a couple of minutes, it probably warrants a comment. Use descriptive variable names and try to follow the [Python style guide](#). At the end of this section, you can find tips and useful commands.

Your code should not take too long to run. For the first few exercises, anything more than a few seconds likely means you are doing something inefficient. Points may be deducted for slow code, even if your solution is correct. You should not need to import any additional libraries. If you think an additional library is necessary, first ask in the forum. It is difficult to name each function you are not allowed to use because especially OpenCV is very powerful and has many of the algorithms we will implement in this course already integrated. As a general guideline, if a task can be solved with one function call, you probably should not use it. If you are unsure, please ask in the TUWEL forum.

Complete the function skeletons provided, implementing edge detection on a grayscale image using the Canny edge detector. The algorithm consists of 4 main parts listed below.

1. Image blurring (2 points)

Blur the image with a Gauss-filter. First, create a 2-dimensional NumPy array which has the values of the Gaussian kernel as defined in the lecture slides. The dimensions of the kernel should be variable depending on `sigma`. The width of the square filter kernel should be calculated as follows:

```
kernel_width = 2 * ceil(3 * sigma) + 1
```

For the actual filtering use the OpenCV function `cv2.filter2D`. Do not use any additional functions of OpenCV. You can use all functions available through NumPy and you should not need to import any additional libraries.

Complete the code in `blur_gauss.py`. For details on the function and the expected return values refer to the docstring in the file.

2. Edge detection (2 points)

Calculate the horizontal and vertical gradient of the image using a Sobel filter, transposing it as appropriate. Again, define the filter kernel as a `np.array` and apply the function `cv2.filter2D`. Make sure that before you do, the image has the data type `np.float32` and is in the range `[0,1]`. Then, calculate the edge strength and the angle of the gradient according to the formulas in the slides. Normalize the edge strengths by dividing by their maximum. If the maximum is zero, return all-zero edge strengths instead of dividing by zero. Do not normalize the values again in the subsequent steps.

Note: You may find `circle.jpg` useful for testing whether your gradient and orientation results are what you expect.

Complete the code in `sobel.py`.

3. Non-maximum suppression (3 points)

Implement the non-maximum suppression algorithm along the orientation of the gradients to thin out the edges. You may wish to quantise the gradient orientations to discrete levels like “horizontal”, “vertical”, “diagonal-left”, “diagonal-right” first and then treat each case separately. E.g. if a pixel has a “horizontal” gradient orientation, its gradient value should only be kept if its immediate left and right neighbor pixels have a lower gradient value than the center pixel. If there are multiple pixels with the same maximum value next to each other, choose one arbitrarily. You can see an example of the image after non-max suppression in Figure 3.

Note: Writing this part of the exercise in a vectorised way (i.e. avoiding loops) is challenging and we do not expect you to do it here. However, it can be done and decreases processing time by more than a factor of 10.

Complete the code in `non_max.py`. Make sure you explain your approach in comments.



Figure 3: Image before and after Non-maximum suppression

4. Hysteresis thresholding (3 points)

Implement hysteresis thresholding to receive connected edge segments as described in the lecture. We recommend you to use the functions `np.where` to find coordinates of pixels in the image with `value ≥ low` and `cv2.connectedComponents` to find connected pixels. After identifying connected pixels, you can use `np.isin` to check each edge if it includes a pixel with `value ≥ high`.

To make parametrisation easier make sure that the gradient values are normalized and lie within the interval $[0, 1]$ before applying the hysteresis thresholding. Perform the threshold comparisons on the floating-point gradient values. For `cv2.connectedComponents`, use an `np.uint8` mask in which background pixels are zero and foreground pixels are nonzero.

Complete the code in `hyst_thresh.py`.

5. Automatic hysteresis threshold selection (2 points)

The hysteresis thresholds can be difficult to select manually, as you will probably discover if you run the code on various images. It's not clear from looking at an image what thresholds you should use. Implement an automatic threshold selection function which calculates these thresholds and then calls the `hyst_thresh` function. The function expects two arguments `low_prop` which is the proportion of edge pixels that should lie above the low threshold, and `high_prop` which is the proportion of edge pixels above the high threshold. This means `low_prop > high_prop`.

Here, edge pixels are the pixels with nonzero gradient values remaining after non-maximum suppression. Exclude zero-valued background pixels when calculating the thresholds. You may use percentiles or sorted values; small differences due to interpolation or rounding are acceptable. Pass the original, spatially arranged gradient image to `hyst_thresh`. If there are no edge pixels, return an all-zero image.

Complete the code in `hyst_auto.py`

Forbidden and allowed commands

Python commands `cv2.Sobel`, `cv2.Canny`, `cv2.GaussianBlur` must not be used. However, the commands `cv2.filter2d` and `np.arctan2`, along with other block-processing and sliding-window functions may be used.

Unless explicitly stated, you should not use functions which solve an exercise, or a significant part of it, in a single line of code. If you are unsure of whether you are allowed to use a function, please contact the tutors.

Advice

Since NumPy is a matrix-oriented library, “vectorised” processing is much more efficient than loops (one or two orders of magnitudes). Try to get familiar with vectorisation as soon as possible. It will decrease processing time significantly which will be necessary especially for later exercises. Tutorials are in the TUWEL course and it is also addressed in the [Python Crash Course video](#).

Images should be converted to floating point format before further processing. Make sure that your starting image only contains RGB values in the uint8 range from 0 to 255. Pay attention to the comments and datatypes in the docstrings of the functions you need to implement. Keep your programs as simple as possible. Catching wrong user inputs and printing error messages is not necessary. The file `helper_functions.py` contains a few functions which might be helpful to plot intermediate results.

Do not modify input arrays in place. Both hysteresis functions should return a binary image with values 0 and 1. If no edges remain, return an all-zero image.

If your code is slow, you can measure which function takes up how much time of the total runtime. Pycharm offers a Profile functionality to do so.

Useful commands

- `cv2.filter2D` - Filtering with linear filter kernel
- `cv2.imread` - Load an image
- `cv2.imshow` - Display an image (for RGB image, value format is uint8 with range [0,255], for grayscale, value format is double with range [0, 1])
- `cv2.waitKey` - Needs to be called after imshow.
- `cv2.cvtColor` - Image color conversion (e.g. to grayscale)
- `cv2.connectedComponents` - Label connected components with a specific label
- `.astype(np.uint8)` - Can be called on np.arrays to convert them to 8-bit unsigned integer format
- `.shape` - Shape of a NumPy array
- `np.amax` - Find maximum
- `np.sort` - Sort an array
- `np.isin` - Return a boolean array of the same shape as the input array with True for elements which are the same in both arrays.
- `np.square` - Elementwise square
- `np.sqrt` - Elementwise square-root
- `np.exp` - Elementwise exponential function with Euler's number as base
- `np.percentile` - Compute the q-th percentile of the data along the specified axis

Part 2: Documentation (6 Points)

Perform the following experiments and answer the questions in a few sentences.

- Answer each numbered question separately, with the answer labelled according to the experiment and question number. For example, Experiment 1 Question 2 should be numbered 1.2. You do not need to include the question text.
- Include images to support your answers. These should be referenced in the answers. The document should include information about what parameter settings were used to generate images so that they can be reproduced. Feel free to use your own images but if you do so, upload the original image as well so we can reproduce your results.
- Your answers should attempt an explanation of why an effect occurs, not just describe what you see in the supporting images.
- Be precise. “It is better” or similarly vague statements without any additional information are not sufficient.

Points may be deducted if your report does not satisfy the above points. Only change the parameters in the file `main.py` in the marked places. We recommend creating a separate Python file for your experiments to not mess up the main file. If you do so, also upload this file because it makes reproducing your experiment results easier for the tutors.

Maximum 1500 words, arbitrary number of pages for result images (referenced in the text). Violating the text limit will lead to point deductions!

Honor Code

Add the honor code to the beginning of the documentation:

For the assignment I followed the honor code:

- Rule 1: Don't look at solutions or code of other students. Everything you submit should be your own work.
- Rule 2: Don't share your solution code with others; however, discussing ideas or general strategies is fine and encouraged.
- Rule 3: Cite the source of any external code directly in your code (e.g., Stack-Overflow). For generative AI tools (e.g., ChatGPT), link to the specific chat. Do not use tools, including coding agents, whose contributions you cannot document and reference.

I know that I have full responsibility for what I submit. I made sure the assignment allows the teaching staff to evaluate which competences I have acquired. I am aware that plagiarism is prohibited and may lead to failure of this course.

Experiment 1 (2 points)

Investigate the effect of `sigma` of the Gaussian blurring.

1. Why is the kernel width defined as a function of sigma? Try various width and sigma inputs for the `cv2.filter2d` function and examine the results. Does an incorrect width have an effect in practice?
2. What effect does blurring have on pixel intensities in the image? You might find the function `plot_row_intensities` from the `helper_functions.py` file useful to document your findings.

Experiment 2 (2 points)

Investigate the effect of thresholds used during hysteresis thresholding.

1. Change the parameters `thresh_low` and `thresh_high`. What effects do these values have on the result?
2. You might notice that keeping thresholds the same over different images can give inconsistent results (try it out). Why is this? Can you think of situations in which you might wish to keep the thresholds the same?

Experiment 3 (2 points)

Investigate the effect of noise on the results. You can add noise to an image with the function `add_gaussian_noise` provided in the `helper_functions.py` file. Add gaussian noise with zero mean and try standard deviations in the range `[0.01, 0.5]`.

1. What effect does the noise have on the resulting edges?
2. Is it possible to set the parameters to get reasonable edges? How does changing `sigma` of `blur_gauss` and the hysteresis threshold values affect the results?

Assistance

Please keep to the following hierarchy whenever there are questions or problems:

1. Forum: Use the TUWEL discussion forum to discuss your problems.
2. Email for questions or in-person help requests: `machinevision@acin.tuwien.ac.at`

Upload

Please upload all Python files and your documentation (pdf-format) **as one ZIP-file** via TUWEL.